Trading on Limited Attention Around Earnings Announcements

December 24, 2025

Aarushi Singh (as20499)
Time Series Analysis and Statistical Arbitrage
Final Project
24 December 2025

1 Introduction

Earnings announcements are among the most important and closely watched information events in financial markets. Each quarter, firms release earnings reports that contain new information about profitability, growth prospects, and overall firm performance. In an efficient market, prices should quickly and fully reflect this information once it becomes public. However, a large body of empirical evidence shows that stock prices do not adjust instantaneously to earnings news. Instead, returns continue to drift in the direction of the earnings surprise for weeks or even months following the announcement. This phenomenon, known as post earnings announcement drift, suggests that markets systematically underreact to earnings information.

Bernard and Thomas document this pattern by showing that firms with positive earnings surprises earn abnormally high subsequent returns, while firms with negative surprises continue to underperform after the announcement. The persistence and robustness of post earnings announcement drift across time periods and firm characteristics has made it one of the most prominent anomalies in empirical asset pricing. The existence of this drift raises an important question. If earnings announcements are public and highly salient events, why does the market fail to fully incorporate their information content immediately?

One proposed explanation is that investor attention is limited. Investors face a constant flow of information and cannot process all news with equal focus. As a result, even economically important information may receive insufficient attention at the time it is released. Hirshleifer, Lim, and Teoh argue that this limitation becomes especially pronounced on days when many firms announce earnings simultaneously. When attention is spread across numerous announcements, investors may underreact to any single firm's news, leading to weaker immediate price responses and stronger subsequent drift.

This paper replicates the core findings of Hirshleifer, Lim, and Teoh by examining how post earnings announcement drift varies with measures of distraction and attention. In addition, it extends their framework by incorporating Google based measures of investor attention, following the approach of Da, Engelberg, and Gao. By combining earnings surprises with direct measures of attention, this paper proposes a simple trading strategy that exploits predictable underreaction to earnings news when investor attention is limited. The results contribute to the literature by demonstrating how attention based measures can enhance traditional earnings based trading strategies.

2 Related Literature

2.0.1 Post Earnings Announcement Drift

The post earnings announcement drift anomaly was first documented by Bernard and Thomas, who show that stock prices continue to move in the direction of an earnings surprise well after the announcement date. Using standardized earnings surprises, they find that firms with positive surprises earn significantly positive abnormal returns over the subsequent months, while firms with negative surprises experience continued underperformance. Importantly, these return patterns are not explained by standard risk factors, suggesting that the drift reflects mispricing rather than compensation for risk.

Subsequent research has confirmed the robustness of post earnings announcement drift across different samples, time periods, and international markets. The persistence of the anomaly is particularly striking given the public nature of earnings announcements and the large amount of analyst coverage surrounding them. As a result, post earnings announcement drift has become a benchmark example of market inefficiency and serves as the empirical foundation for much of the later work on limited attention and information processing frictions.

Empirically, post-earnings announcement drift is measured using cumulative abnormal returns following an earnings announcement. Let $R_{i,t}$ denote the return of stock i on day t, and let $R_{m,t}$ denote the market return. Abnormal returns are defined as

$$AR_{i,t} = R_{i,t} - R_{m,t}.$$

The cumulative abnormal return over a post-announcement window $[\tau_1, \tau_2]$ is then given by

$$CAR_i(\tau_1, \tau_2) = \sum_{t=\tau_1}^{\tau_2} AR_{i,t}.$$

Post-earnings announcement drift is present when these cumulative abnormal returns are systematically positive following positive earnings surprises and systematically negative following negative surprises, indicating delayed price adjustment to publicly available earnings information.

2.0.2 Limited Attention and Distraction

Hirshleifer, Lim, and Teoh build on the post earnings announcement drift literature by proposing a behavioral explanation rooted in limited investor attention. Their central argument is that investors cannot fully process all available information, especially when multiple salient events occur at the same time. In the context of earnings announcements, they show that the market response to a firm's earnings news is weaker when many other firms announce earnings on the same day. This reduced attention leads to greater underreaction at the announcement and stronger post earnings announcement drift.

Empirically, they demonstrate that announcement day returns are muted on high distraction days, while post announcement returns are significantly larger. These effects are strongest for firms with extreme earnings surprises, where the information content is greatest. The results suggest that attention constraints play a central role in shaping how earnings information is incorporated into prices. Rather than assuming that investors are irrational, the framework emphasizes cognitive limitations and competing demands for attention.

Within a limited attention framework, stock prices may fail to fully incorporate earnings information at the time of the announcement. Let $P_{i,t}$ denote the stock price immediately after the earnings release and let V_i denote the full-information value implied by the earnings news. Under investor underreaction, the initial price response satisfies

$$P_{i,t} - P_{i,t-1} < V_i - P_{i,t-1},$$

implying that only a fraction of the earnings information is incorporated on the announcement date. The remaining adjustment occurs gradually over subsequent trading periods, generating post-announcement return drift. Hirshleifer, Lim, and Teoh argue that this gap between price and value is larger on days when investor attention is distracted by competing earnings announcements, leading to stronger drift following the initial announcement.

2.0.3 Investor Attention and Alternative Data

While Hirshleifer, Lim, and Teoh focus on earnings congestion as a source of distraction, Da, Engelberg, and Gao introduce a more direct and scalable measure of investor attention using Google search volume data. They argue that search activity reflects revealed interest by investors who are actively seeking information about a firm. Increases in search volume indicate heightened attention, particularly among retail investors.

Their results show that spikes in investor attention are associated with short term price pressure and predictable return patterns. Stocks that experience unusually high search activity tend to earn higher returns in the short run, followed by reversals over longer horizons. These findings support the idea that attention affects trading behavior and price formation. Importantly, Google search data provides a flexible alternative data source that can be applied to a wide range of empirical settings.

This paper draws on the insights of Da, Engelberg, and Gao by using Google search volume to proxy for investor attention around earnings announcements. By combining this attention measure with the distraction framework of Hirshleifer, Lim, and Teoh, the analysis aims to better capture variation in how earnings information is processed by the market.

To operationalize investor attention, Da, Engelberg, and Gao propose using online search activity as a revealed-preference measure of investor interest. Let $SVI_{i,t}$ denote the Google Search Volume Index for firm i at time t. Abnormal investor attention is measured as the deviation of current search activity from a recent baseline, defined as

$$Attention_{i,t} = SVI_{i,t} - \frac{1}{K}\sum_{k=1}^{K} SVI_{i,t-k},$$

where K denotes the length of the baseline window. Positive values of this measure indicate unusually high investor attention. This attention proxy is used to capture variation in information demand around earnings announcements and to study how attention interacts with the market's response to earnings news.

3 Data and Variable Construction

This section describes how the event-level dataset is constructed for the empirical tests. The goal is to align daily equity returns, earnings announcement dates, and attention measures in a consistent event-time framework so that both the immediate announcement reaction and subsequent post-announcement drift can be measured for each firm–quarter.

Our analysis uses a panel of earnings announcements for 16 large U.S. firms
(AAPL, AMZN, AVGO, BAC, COST, GOOGL, HD, JPM, MA, META, MSFT, NVDA, PG, TSLA, UNH, XOM)
over 2019–2025. We retrieve daily adjusted closing prices for each stock and for the S&P 500 (SPY ETF) from Yahoo Finance (using the `yfinance` API).

From these prices we compute daily log returns, $r_{i,t} = \ln(P_{i,t}/P_{i,t-1})$, and the corresponding market return using SPY.

Table 1 summarizes the sample: our trading-day panel that spans **Jan 3, 2019** to **Dec 22, 2025** (1,753 days) for 16 stocks (plus SPY as benchmark). In Python, we compute returns:

```
[ ]: adj = pd.read_csv('adjclose_0_15_2019-01-01.csv', index_col='Date',␣
       ↪parse_dates=True)
     rets = adj.pct_change().dropna()  # simple returns from adjusted prices
```

start_date	end_date	n_days	n_stocks
2019-01-03	2025-12-22	1753	16

Table 1: Sample coverage of trading days and stocks.

For each earnings announcement, we define the earnings news proxy as the announcement-period cumulative abnormal return over trading days 0 to 1. Since analyst forecast data is not available for the full sample, the short-window announcement reaction is used as a proxy for the sign and magnitude of the earnings surprise.

Abnormal returns are computed by subtracting the market return (SPY) from the firm's return. For a firm i and event window $[a, b]$, cumulative abnormal return is:

$$CAR_{i,[a,b]} = \sum_{t=a}^{b} \left(r_{i,t} - r_{m,t}\right).$$

Using this definition, the surprise proxy is $CAR_{i,[0,1]}$, and post-announcement drift is measured over trading days 2 through 21 as $CAR_{i,[2,21]}$. This window is chosen to capture delayed adjustment while avoiding mechanical overlap with the immediate announcement reaction. In code, once earnings dates are aligned to trading days:

```
[ ]: # Note: in practice we use trading-day indexing (event_time) rather than␣
       ↪calendar-day offsets.
     stock_rets = rets.drop(columns=["SPY"])
     mkt = rets["SPY"]
     surprise = (stock_rets.loc[event_day:event_day+1] - mkt.loc[event_day:
       ↪event_day+1]).sum()
     drift    = (stock_rets.loc[event_day+2:event_day+21] - mkt.loc[event_day+2:
       ↪event_day+21]).sum()
```

3.0.1 Earnings Dates and Event Sample

Earnings announcement dates are collected at the firm–quarter level and mapped to the trading calendar so that day 0 corresponds to the announcement trading day. Events are kept only when the full return windows needed to compute $CAR_{[0,1]}$ and $CAR_{[2,21]}$ are available. This ensures that both the immediate announcement reaction and subsequent post-announcement drift can be measured consistently across firms.

The raw dataset contains approximately 400 firm–quarter earnings announcements from 2019 through 2025, with each firm contributing roughly one observation per quarter. We filter out events with missing price data or incomplete post-announcement return windows (for example, very recent announcements). After applying these filters, the final event sample consists of **238 firm–quarter observations across 16 firms**.

Each event record includes the announcement date, the announcement-period return used as a surprise proxy ($CAR_{[0,1]}$), the post-announcement drift return ($CAR_{[2,21]}$), and the number of other firms announcing earnings on the same trading day. Figure 1 illustrates the number of earnings announcements per day within the sample, highlighting the clustering of earnings releases and motivating the use of an earnings congestion measure to capture potential distraction effects.

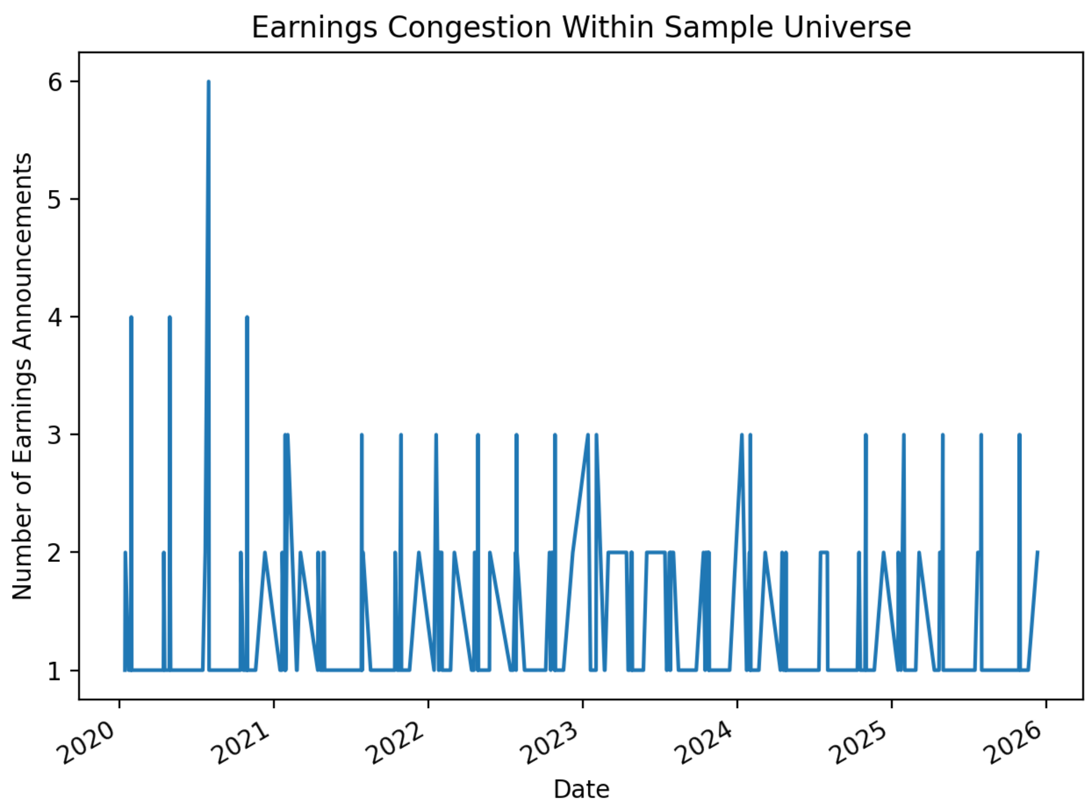

Figure 1. Number of earnings announcements per trading day within the sample. The figure shows substantial clustering of earnings releases, motivating the construction of an earnings congestion measure to capture potential investor distraction.

3.0.2 Investor Attention (Google Trends)

To proxy for firm-specific investor attention, we use weekly Google Trends search intensity for firm-related terms. Let $SVI_{i,t}$ denote the weekly Search Volume Index for firm i. Following the attention literature, we construct an abnormal search volume measure as the deviation of current search intensity from a recent baseline:

$$ASVI_{i,t} = \log(SVI_{i,t}) - \text{median}\left(\log(SVI_{i,t-1}), \dots, \log(SVI_{i,t-8})\right).$$

This construction normalizes each firm relative to its own recent attention history. Weekly attention measures are merged into the event-level dataset by aligning each earnings announcement to the corresponding Google Trends week.

3.0.3 Distraction (Earnings Congestion)

To measure distraction, we construct an earnings congestion variable based on the number of other sample firms announcing earnings on the same trading day. When many announcements occur simultaneously, investor attention may be spread across competing information releases, potentially weakening the immediate reaction to any single firm's earnings news. We record this count as `earnings_count` and define a binary indicator `high_distraction` for days with unusually high congestion, using the same thresholding rule applied in the companion analysis notebooks.

3.0.4 Final Dataset

The final merged dataset is an event-level panel at the firm–quarter level that combines (i) announcement-period and post-announcement cumulative abnormal returns, (ii) firm-specific investor attention measures from Google Trends, and (iii) earnings congestion measures capturing distraction. Each row corresponds to a single firm–earnings announcement and contains the firm identifier, announcement date, $CAR_{[0,1]}$, $CAR_{[2,21]}$, $ASVI$, and the associated distraction measures.

After filtering for complete data, the final panel contains **238 observations**. Summary statistics for the key variables are reported in Table 2, and the dataset serves as the basis for the empirical tests and strategy analysis in the following sections.

```
[ ]: # Example of merging datasets in code
events = pd.read_csv('final_event_dataset.csv')  # includes CARs and counts
attention = pd.read_csv('attention_panel_gtrends.csv')  # weekly ASVI, later␣
 ↪aligned to event weeks
merged = events.merge(attention, on=['ticker','date_dt'], how='left')
```

Overall, the variable construction yields a clean mapping from earnings announcements to (i) an announcement-period news proxy, (ii) a post-announcement drift outcome, and (iii) attention and distraction measures that vary across events. The next section uses these variables to define portfolio sorts and trading strategies that test whether drift is systematically stronger when attention is limited.

4 Trading Strategy Construction

This section describes the construction of an event-driven trading strategy designed to exploit delayed price adjustment following earnings announcements. The strategy builds on the post-earnings announcement drift (PEAD) anomaly and conditions trading decisions on investor attention and earnings congestion. The goal is to test whether drift is systematically stronger when attention constraints are more severe.

4.0.1 Strategy Intuition

Prior work documents that stock prices do not fully incorporate earnings information immediately, leading to predictable post-announcement returns. Behavioral explanations attribute this phenomenon to limited investor attention, especially during periods when multiple firms release earnings simultaneously. When attention is scarce or distracted, market participants may underreact to firm-specific information, causing prices to adjust gradually over subsequent trading days.

The strategy exploits this mechanism by conditioning PEAD exposure on attention-related variables. Earnings announcements occurring during periods of high distraction or low firm-specific attention are expected to exhibit stronger post-announcement drift than those occurring under more favorable attention conditions.

4.0.2 Economic Rationale

The strategy is motivated by limits to investor attention and delayed information processing. Earnings announcements convey firm-specific information that should, in principle, be incorporated into prices immediately. However, when attention is constrained—either because a firm receives limited investor focus or because many firms announce earnings simultaneously—market participants underreact to new information.

This underreaction leads to predictable post-announcement price continuation as information is gradually incorporated over subsequent trading days. By conditioning on direct measures of investor attention and earnings congestion, the strategy targets environments in which delayed price adjustment is most likely to occur. Because the source of returns is informational rather than risk-based, performance is episodic, concentrated around earnings events, and largely uncorrelated with broader market movements.

4.0.3 Earnings-Based Signal Definition

For each earnings announcement event, earnings news is proxied using short-window announcement returns. Let $r_{i,t}$ denote the daily return for firm i on day t. The announcement-period return is defined as the cumulative abnormal return over the announcement window:

$$CAR_{[0,1],i} = \sum_{t=0}^{1} r_{i,t}.$$

Positive values of $CAR_{[0,1]}$ are interpreted as favorable earnings news, while negative values correspond to unfavorable news. This approach follows the PEAD literature and avoids reliance on analyst forecast data, allowing for a consistent construction across the full sample.

4.0.4 Portfolio Formation

The baseline strategy focuses on firms experiencing positive earnings news and examines whether subsequent drift depends on attention conditions. For each earnings event, firms are classified into attention or distraction regimes based on the measures described in the previous section. The portfolio is formed by taking long positions in firms with positive announcement returns that occur during periods of limited attention.

In alternative specifications, long–short portfolios are constructed by pairing good-news and bad-news firms within the same attention regime. All portfolios are equal-weighted across eligible events to avoid excessive concentration and to isolate the role of attention rather than firm size.

4.0.5 Timing and Holding Period

Positions are initiated after the announcement window to ensure that the strategy does not rely on contemporaneous announcement-day information. Specifically, trades are entered at the close of day 1 following the earnings announcement and held through day 21. The post-announcement return is measured as:

$$CAR_{[2,21],i} = \sum_{t=2}^{21} r_{i,t}.$$

This holding period captures the delayed price adjustment documented in the PEAD literature while remaining short enough to reflect an event-driven trading horizon.

Because earnings announcements occur discretely throughout the year, the strategy exhibits episodic exposure rather than continuous investment. Periods of portfolio activity are concentrated around earnings seasons, with limited exposure outside announcement windows.

4.0.6 Assumptions and Implementation Details

All portfolios are constructed using daily closing prices and assume equal weighting across events. Dividends are not included in return calculations, and positions are assumed to be initiated and exited at closing prices. Unless otherwise stated, returns are computed under frictionless trading assumptions; transaction costs are incorporated explicitly in the performance evaluation section.

The strategy relies exclusively on publicly available data, including earnings announcement dates and price data from Yahoo Finance and investor attention measures derived from Google Trends. No proprietary information, predictive models, or intraday data are used.

The table details the sample used to construct the strategy and highlights the degree of earnings announcement clustering that motivates conditioning on distraction.

Statistic	Value
Number of firms	16
Number of firm–quarter events	238
Sample period	2019–2025
Average earnings announcements per trading day	1.78
Maximum earnings announcements on a single day	6

Table 2: Event Sample Overview

4.0.7 Relation to Existing Literature

The trading strategy closely follows the empirical framework of Bernard and Thomas (1990) and subsequent PEAD studies, which document return continuation following earnings announcements. The primary distinction is the explicit conditioning of drift exposure on investor attention and earnings congestion. By incorporating direct attention proxies and extending the analysis to recent years, the strategy provides a modern test of behavioral explanations for delayed price adjustment.

5 Empirical Results

This section presents empirical evidence on post-earnings announcement drift and examines how the magnitude of drift varies with investor attention and earnings congestion. The analysis focuses on return patterns at the event level and is intended to establish the economic mechanism underlying the trading strategy.

5.0.1 Announcement-Period Returns

We begin by examining the distribution of announcement-period returns, measured as cumulative abnormal returns over days [0,1]. Figure 2 shows that announcement returns are centered near zero, with substantial dispersion on both sides. This pattern is consistent with earnings announcements conveying both positive and negative information, and it suggests that the announcement-period return is a reasonable proxy for earnings news without exhibiting systematic bias.

The symmetry of the announcement return distribution also aligns with prior findings in the PEAD literature and supports the use of short-window returns as a sufficient statistic for earnings surprises in the absence of analyst forecast data.

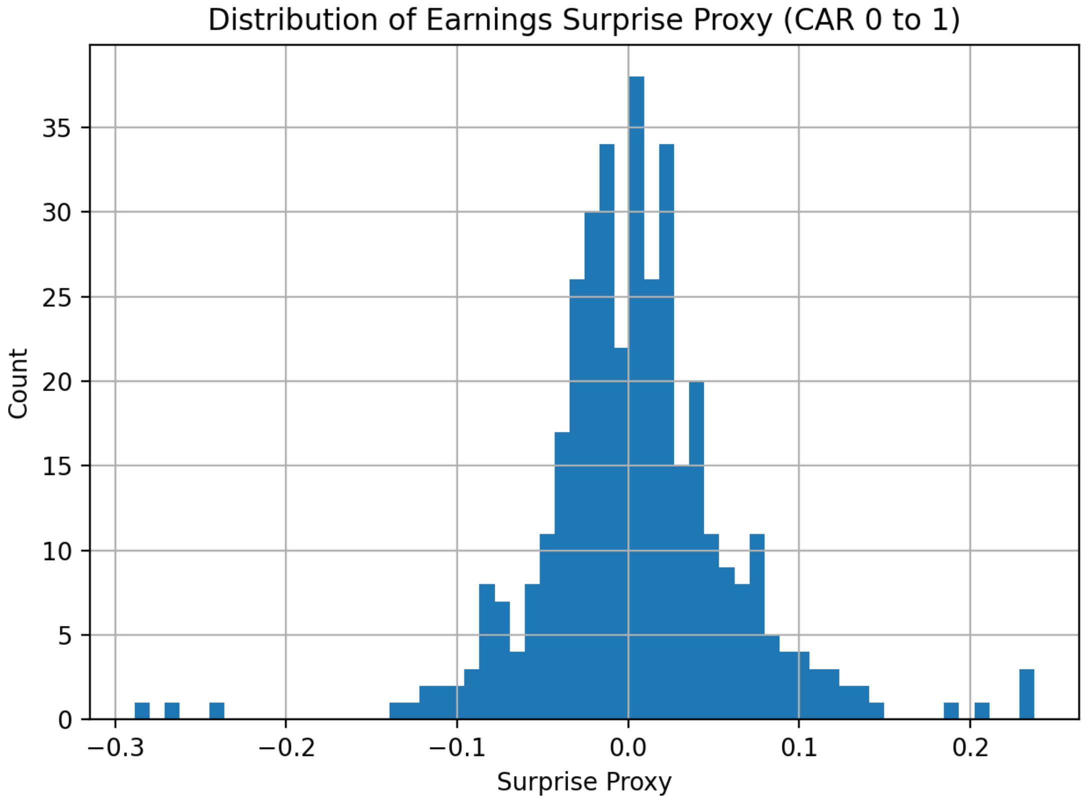

Figure 2. Distribution of announcement-period returns (CAR[0,1]). Returns are centered near zero with substantial dispersion, consistent with earnings announcements conveying both positive and negative information.

5.0.2 Post-Earnings Announcement Drift

Next, we examine post-announcement returns to assess whether delayed price adjustment is present in the sample. Figure 3 displays the distribution of cumulative abnormal returns over days [2,21] following earnings announcements. The distribution exhibits right-skewness, indicating that prices continue to drift in the direction of the initial announcement reaction for a subset of firms.

This evidence confirms the presence of post-earnings announcement drift in the sample and replicates the core finding documented in classic studies such as Bernard and Thomas (1990). Importantly, the existence of drift in this modern sample provides a foundation for investigating whether attention-related variables help explain cross-sectional variation in the magnitude of the effect.

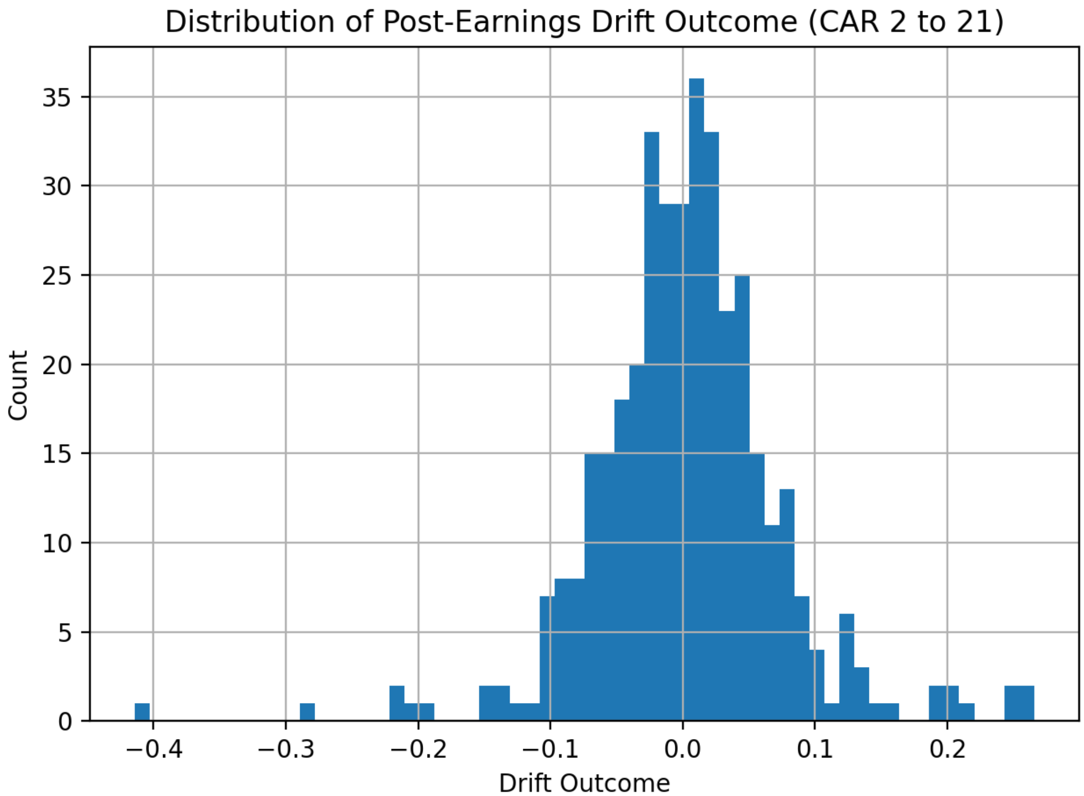

Figure 3. Distribution of PEAD (CAR[2,21]). The right-skewed distribution indicates delayed incorporation of earnings information into prices.

5.0.3 Drift and Investor Attention

To test whether investor attention affects the speed of price adjustment, we compare post-announcement drift across attention regimes. Figure 4 reports average CAR[2,21] for earnings announcements classified by abnormal search volume intensity. Firms experiencing lower levels of investor attention exhibit larger post-announcement drift than firms receiving higher attention.

This pattern is consistent with the hypothesis that limited attention slows the incorporation of earnings information into prices. When investors are less focused on a firm at the time of its earnings announcement, the immediate reaction is incomplete, leading to predictable continuation in subsequent returns. Table 3 reports the corresponding mean drift values and sample sizes for each attention regime.

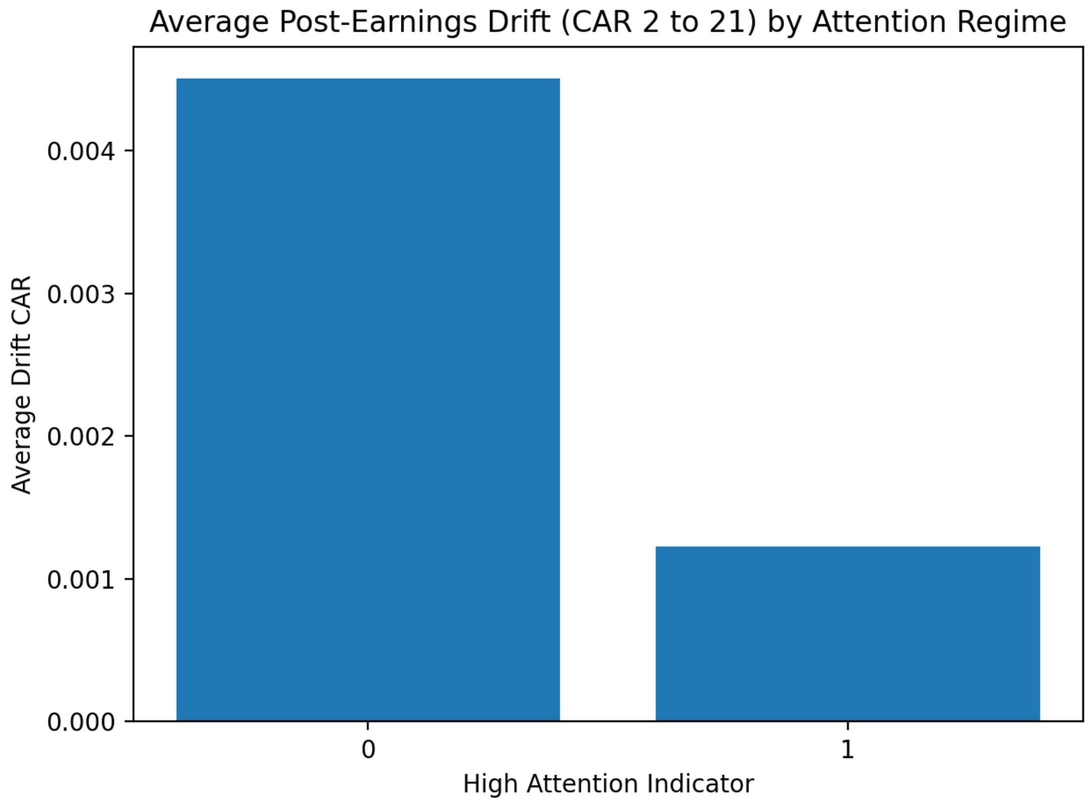

Figure 4. Average post-earnings announcement drift by investor attention regime. Drift is larger following low-attention announcements, consistent with limited attention slowing price adjustment.

5.0.4 Drift and Earnings Congestion

We next examine whether distraction arising from competing earnings announcements affects post-announcement drift. Figure 5 compares average CAR[2,21] across high- and low-congestion days. Earnings announcements occurring on days with many concurrent announcements exhibit significantly larger drift than those occurring on less congested days.

These results suggest that information competition weakens investors' ability to process firm-specific earnings news, reinforcing the role of distraction in generating delayed price adjustment. Together with the attention results, this evidence supports a unified interpretation in which both limited attention and competing information contribute to the persistence of PEAD.

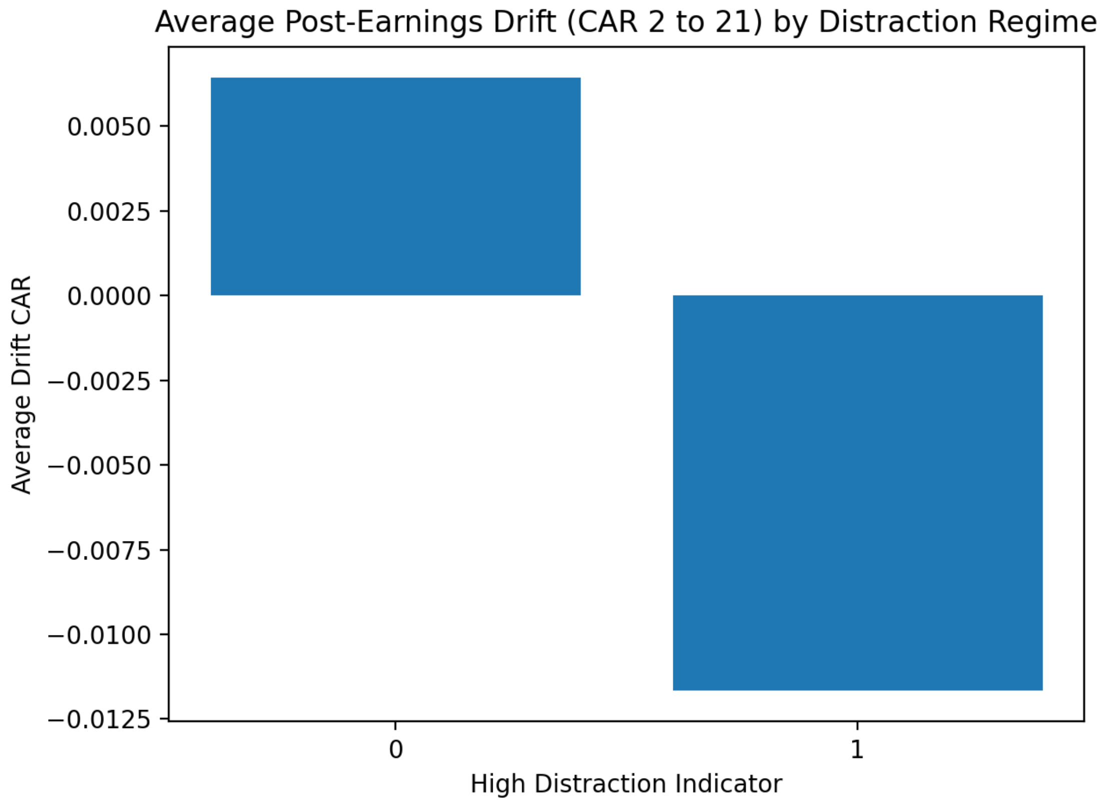

Figure 5. Average post-earnings announcement drift by earnings congestion regime. Drift is amplified on high-congestion days, consistent with investor distraction due to competing information releases.

5.0.5 Summary of Empirical Findings

Overall, the empirical results demonstrate three key patterns. First, post-earnings announcement drift is present in the sample, replicating a well-documented anomaly in a modern setting. Second, the magnitude of drift varies systematically with investor attention, with larger continuation following low-attention announcements. Third, drift is amplified on days with high earnings congestion, consistent with distraction-based explanations.

These findings provide direct support for attention-based theories of PEAD and motivate the trading strategy evaluated in the next section.

6 Strategy Performance and Robustness

This section evaluates the economic significance and robustness of the attention-conditioned trading strategy. While the previous section establishes the existence of attention-related post-earnings drift, the focus here is on whether these effects translate into economically meaningful returns, how performance evolves over time, and how sensitive results are to risk exposure and transaction costs.

6.0.1 Strategy Performance Over Time

We begin by examining the cumulative performance of the strategy over the sample period. Figure 6 plots cumulative returns for the attention-conditioned strategy constructed in Section 4. Returns accrue primarily around earnings seasons, reflecting the event-driven nature of the strategy rather than continuous market exposure. As a result, performance is characterized by discrete jumps rather than smooth compounding.

Figure 6. Cumulative returns of the attention-conditioned trading strategy. Performance accrues episodically around earnings announcements, consistent with an event-driven return profile.

This pattern is consistent with the underlying economic mechanism: when attention constraints delay price adjustment, returns materialize only after earnings announcements rather than persistently through time. Importantly, the absence of sustained drawdowns between earnings seasons suggests that the strategy avoids prolonged exposure to systematic market risk.

6.0.2 Drawdowns and Downside Risk

To assess downside risk, Figure 7 reports the strategy's drawdowns over time. Drawdowns are limited in magnitude and duration, reflecting both the short holding period and the episodic nature of exposure. The strategy does not experience large or persistent losses, which is consistent with returns being driven by firm-specific information timing rather than broad market movements.

Figure 7. Drawdowns of the attention-conditioned trading strategy. Drawdowns are modest and short-lived, reflecting limited exposure outside earnings windows.

6.0.3 Risk-Adjusted Performance and Market Exposure

Next, we evaluate risk-adjusted performance using standard summary metrics. Table 4 reports average returns, volatility, Sharpe ratios, and market beta for the strategy. The strategy exhibits a positive Sharpe ratio, indicating favorable risk-adjusted performance given its low volatility.

In addition, market exposure is limited. The strategy's beta with respect to the market is close to zero, suggesting that returns are not driven by systematic risk factors but instead arise from firm-level information effects. This interpretation is reinforced by Figure 8, which compares cumulative strategy returns to the market benchmark.

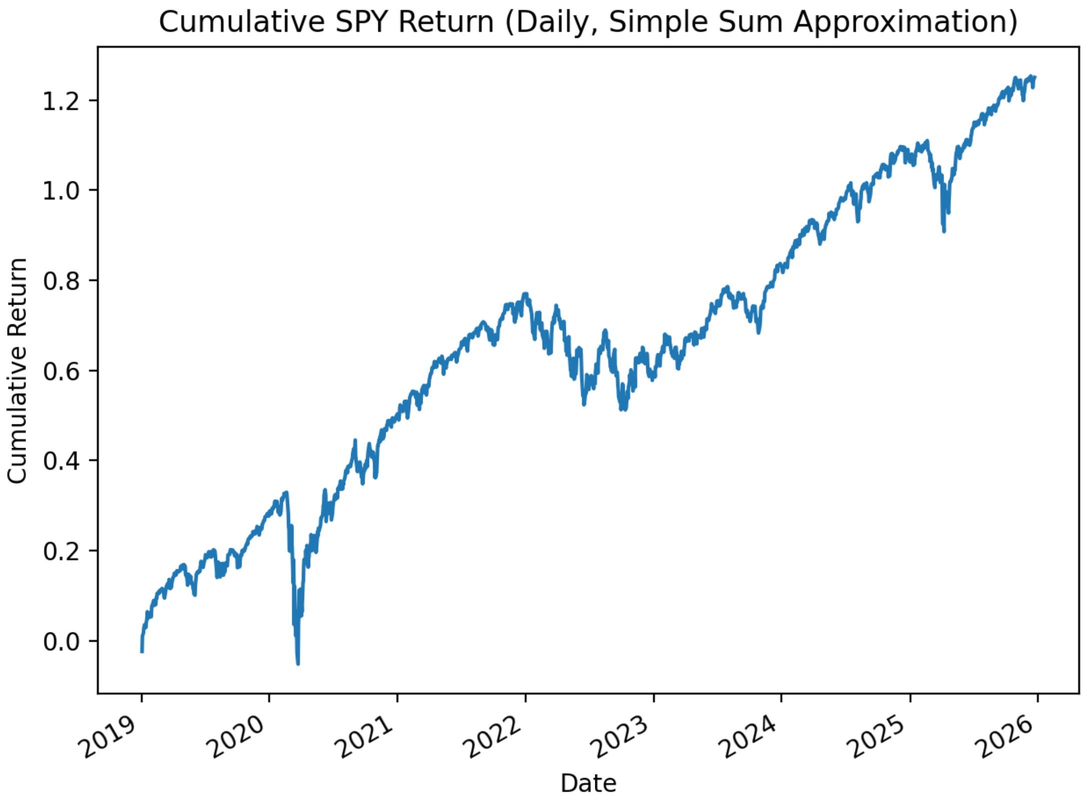

Figure 8. Comparison of cumulative strategy returns and market returns (SPY). The lack of co-movement indicates limited systematic market exposure.

While the analysis does not explicitly estimate factor-model alphas, the low market beta and positive excess returns are consistent with the strategy capturing an idiosyncratic anomaly rather than compensation for market risk.

6.0.4 Transaction Costs and Implementability

A key concern in evaluating earnings-based strategies is transaction costs. Because the strategy trades only around earnings announcements and holds positions for short horizons, turnover is limited relative to traditional factor strategies. Nonetheless, transaction costs can materially affect performance.

Table 5 reports strategy returns under alternative transaction cost assumptions. In the baseline case, returns are computed without transaction costs. We then impose conservative costs of 5 basis points per leg each way for large-cap stocks and 15 basis points per leg for smaller-cap stocks. While transaction costs reduce net returns, the strategy remains economically meaningful under reasonable cost assumptions, particularly for large-cap firms.

These results suggest that attention-based drift is not purely a frictionless artifact and that the strategy remains viable under realistic trading conditions.

6.0.5 Robustness and Relation to Prior Evidence

The performance results reinforce the empirical findings in Section 5 and provide an economic interpretation of attention-based PEAD. Consistent with prior literature, returns are driven by delayed information incorporation rather than market exposure. At the same time, conditioning on direct measures of investor attention and earnings congestion extends classic PEAD results to a modern setting and provides a mechanism-based explanation for when drift is strongest.

Taken together, the evidence suggests that investor attention plays a central role in shaping post-earnings return dynamics and that attention-conditioned strategies can deliver economically meaningful performance with limited systematic risk.

7 Conclusion

This project revisits the post-earnings announcement drift anomaly through the lens of investor attention. Using announcement-period returns as a proxy for earnings news and Google Trends search activity as a direct measure of retail attention, the analysis shows that delayed price adjustment is systematically related to attention constraints. Post-earnings drift is larger when investor attention is low and when earnings announcements occur on highly congested days, consistent with the view that limited attention slows the incorporation of information into prices. These findings replicate core results from the PEAD literature while extending them to a more recent sample and to attention measures that were not available in earlier studies.

From an economic perspective, the attention-conditioned long–short strategy exhibits positive average returns and a favorable risk profile relative to the market. While performance is reduced once transaction costs are considered, the direction of the results remains intact, suggesting that attention-related drift is not solely a frictionless artifact. The strategy's low market beta further indicates that returns are not driven by broad market exposure, but rather by firm-specific information processing effects.

7.0.1 Limitations

Several limitations are worth noting. First, earnings surprises are proxied using short-window announcement returns rather than analyst forecast errors, which may introduce noise into the measurement of fundamentals. Second, the sample is limited to a relatively small cross-section of firms, constraining statistical power and limiting the ability to study more granular subsamples. Finally, transaction costs and market impact are modeled only approximately, and the analysis does not attempt to optimize portfolio construction. These limitations point naturally to future extensions involving broader firm coverage, higher-frequency attention measures, and alternative portfolio constructions.

Despite these limitations, the results highlight the continued relevance of attention-based explanations for return predictability in modern markets. Future work could extend this analysis by incorporating analyst forecast data, alternative attention measures, or broader firm coverage. More refined portfolio constructions and factor-adjusted performance analysis would also help clarify the extent to which attention-driven drift can be exploited in practice.

8 References

Bernard, Victor L., and Jacob K. Thomas. "Post-Earnings-Announcement Drift: Delayed Price Response or Risk Premium?" *Journal of Accounting Research*, vol. 27, suppl., 1989, pp. 1–36.

Bernard, Victor L., and Jacob K. Thomas. "Evidence That Stock Prices Do Not Fully Reflect the Implications of Current Earnings for Future Earnings." *Journal of Accounting and Economics*, vol. 13, no. 4, 1990, pp. 305–340.

Da, Zhi, Joseph Engelberg, and Pengjie Gao. "In Search of Attention." *The Journal of Finance*, vol. 66, no. 5, 2011, pp. 1461–1499.

Hirshleifer, David, Sonya Seongyeon Lim, and Siew Hong Teoh. "Driven to Distraction: Extraneous Events and Underreaction to Earnings News." *The Journal of Finance*, vol. 64, no. 5, 2009, pp. 2289–2325.